

```
#Change runtime type
#Google Colab: Use the menu option "Runtime > Change Runtime Type" and select "GPU" from the "Hardware Accelerator" dropdown.
#The code is run in google server which give us some free GPU with limited access
#windows and linux: NVIDIA graphics card ---> compatibles NVIDIA cuda driver
#binder and MacOs
#kaggle - setting - GPU - Accelerator
```

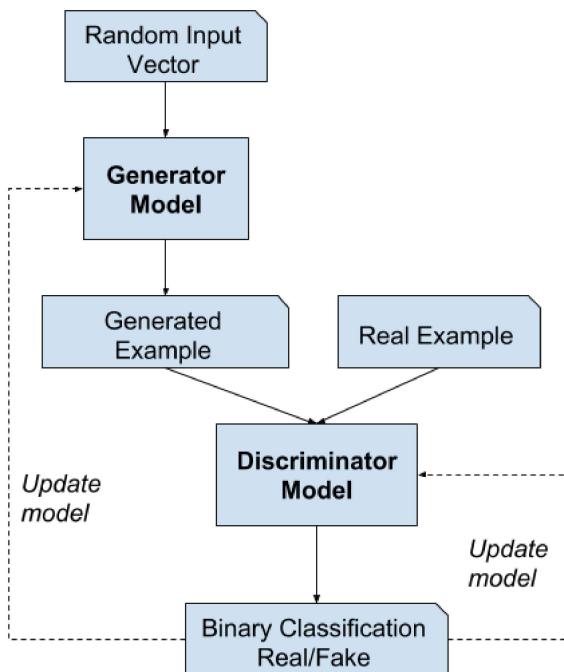
❖ Introduction to Generative Modeling

Deep neural networks are used mainly for supervised learning: classification or regression. Generative Adversarial Networks or GANs, however, use neural networks for a very different purpose: Generative modeling

Generative modeling is an unsupervised learning task in machine learning that involves automatically discovering and learning the regularities or patterns in input data in such a way that the model can be used to generate or output new examples that plausibly could have been drawn from the original dataset. - [Source](#)

To get a sense of the power of generative models, just visit thispersondoesnotexist.com. Every time you reload the page, a new image of a person's face is generated on the fly. The results are pretty fascinating:

While there are many approaches used for generative modeling, a Generative Adversarial Network takes the following approach:



There are two neural networks: a *Generator* and a *Discriminator*. The generator generates a "fake" sample given a random vector/matrix, and the discriminator attempts to detect whether a given sample is "real" (picked from the training data) or "fake" (generated by the generator). Training happens in tandem: we train the discriminator for a few epochs, then train the generator for a few epochs, and repeat. This way both the generator and the discriminator get better at doing their jobs.

GANs however, can be notoriously difficult to train, and are extremely sensitive to hyperparameters, activation functions and regularization. In this tutorial, we'll train a GAN to generate images of anime characters' faces.



```
project_name = 'NU Project '
```

▼ Downloading and Exploring the Data

We can use the [opendatasets](#) library to download the [dataset](#) from Kaggle. opendatasets uses the [Kaggle Official API](#) for downloading datasets from Kaggle. Follow these steps to find your API credentials:

1. Sign in to <https://kaggle.com/>, then click on your profile picture on the top right and select "My Account" from the menu.
2. Scroll down to the "API" section and click "Create New API Token". This will download a file kaggle.json with the following contents:

```
{"username": "YOUR_KAGGLE_USERNAME", "key": "YOUR_KAGGLE_KEY"}
```

3. When you run `opendatasets.download`, you will be asked to enter your username & Kaggle API, which you can get from the file downloaded in step 2.

Note that you need to download the `kaggle.json` file only once. On Google Colab, you can also upload the `kaggle.json` file using the files tab, and the credentials will be read automatically.

```
!pip install opendatasets --upgrade --quiet
#Importing opendataset using pip

import opendatasets as od
dataset_url = 'https://www.kaggle.com/datasets/splcher/animefacedataset'
od.download(dataset_url)

#key: ishikataklikar
#password:9940351ceafcd125115c6a9c505ac4d2a
#63000+ img
#The dataset has a single folder called images which contains all 63,000+ images in JPG format.

→ Downloading animefacedataset.zip to ./animefacedataset
100% [██████████] 395M/395M [00:13<00:00, 31.8MB/s]
```

```
import os
DATA_DIR = './animefacedataset'
print(os.listdir(DATA_DIR))

→ ['images']

print(os.listdir(DATA_DIR+'/images')[:10])

→ ['12404_2005.jpg', '42341_2013.jpg', '22056_2008.jpg', '3849_2002.jpg', '34176_2011.jpg', '42162_2013.jpg', '7678_2004.jpg', '3170_2011.jpg', '22056_2008.jpg', '3849_2002.jpg']
```

Let's load this dataset using the `ImageFolder` class from `torchvision`. We will also resize and crop the images to 64x64 px, and normalize the pixel values with a mean & standard deviation of 0.5 for each channel. This will ensure that pixel values are in the range (-1, 1), which is more convenient for training the discriminator. We will also create a data loader to load the data in batches.

```
from torch.utils.data import DataLoader
from torchvision.datasets import ImageFolder
import torchvision.transforms as T

image_size = 64
batch_size = 128 #for creating a data loader
stats = (0.5, 0.5, 0.5), (0.5, 0.5, 0.5)
#mean # SD

train_ds = ImageFolder(DATA_DIR, transform=T.Compose([
    T.Resize(image_size),
    T.CenterCrop(image_size),
    T.ToTensor(),
    T.Normalize(*stats)])) #T refer to torch vision transform and compose is for the batch of images #64*64

train_dl = DataLoader(train_ds,
                     batch_size,
                     shuffle=True,
```

```
    num_workers=3,
    pin_memory=True)

→ /usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.py:557: UserWarning: This DataLoader will create 3 worker processes.
  warnings.warn(_create_warning_msg)
```

Let's create helper functions to denormalize the image tensors and display some sample images from a training batch.

```
import torch
from torchvision.utils import make_grid
import matplotlib.pyplot as plt
%matplotlib inline

# lets visualize the training dataset

def denorm(img_tensors):
    return img_tensors * stats[1][0] + stats[0][0]

def show_images(images, nmax=64):
    fig, ax = plt.subplots(figsize=(8, 8))
    ax.set_xticks([]); ax.set_yticks([])
    ax.imshow(make_grid(denorm(images.detach()[:nmax]), nrow=8).permute(1, 2, 0))

def show_batch(dl, nmax=64):
    for images, _ in dl:
        show_images(images, nmax)
        break

show_batch(train_dl)
```



```
def get_default_device():
    """Pick GPU if available, else CPU"""
    if torch.cuda.is_available():
        return torch.device('cuda')
    else:
        return torch.device('cpu')

def to_device(data, device):
    """Move tensor(s) to chosen device"""
    if isinstance(data, (list, tuple)):
        return [to_device(x, device) for x in data]
    return data.to(device, non_blocking=True)
```

```

class DeviceDataLoader():
    """Wrap a dataloader to move data to a device"""
    def __init__(self, dl, device):
        self.dl = dl
        self.device = device

    def __iter__(self):
        """Yield a batch of data after moving it to device"""
        for b in self.dl:
            yield to_device(b, self.device)

    def __len__(self):
        """Number of batches"""
        return len(self.dl)

device = get_default_device()
device

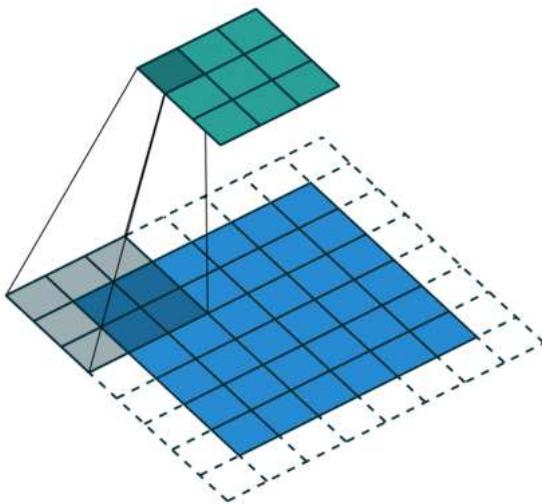
device(type='cuda')

train_dl = DeviceDataLoader(train_dl, device)

```

▼ Discriminator Network

The discriminator takes an image as input, and tries to classify it as "real" or "generated". In this sense, it's like any other neural network. We'll use a convolutional neural networks (CNN) which outputs a single number output for every image. We'll use stride of 2 to progressively reduce the size of the output feature map.



```

import torch.nn as nn

discriminator = nn.Sequential(
    # in: 3 x 64 x 64
    nn.Conv2d(3, 64, kernel_size=4, stride=2, padding=1, bias=False),
    nn.BatchNorm2d(64),
    nn.LeakyReLU(0.2, inplace=True),
    # out: 64 x 32 x 32

    nn.Conv2d(64, 128, kernel_size=4, stride=2, padding=1, bias=False),
    nn.BatchNorm2d(128),
    nn.LeakyReLU(0.2, inplace=True),
    # out: 128 x 16 x 16

    nn.Conv2d(128, 256, kernel_size=4, stride=2, padding=1, bias=False),
    nn.BatchNorm2d(256),
    nn.LeakyReLU(0.2, inplace=True),
    # out: 256 x 8 x 8

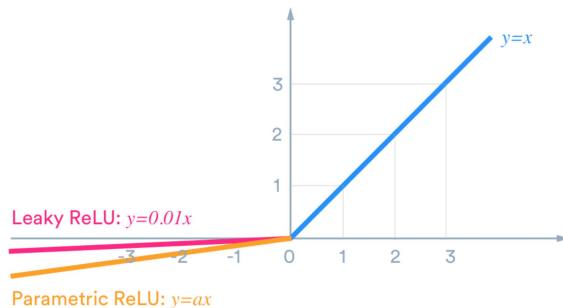
    nn.Conv2d(256, 512, kernel_size=4, stride=2, padding=1, bias=False),
    nn.BatchNorm2d(512),
    nn.LeakyReLU(0.2, inplace=True),
    # out: 512 x 4 x 4

    nn.Conv2d(512, 1, kernel_size=4, stride=1, padding=0, bias=False),

```

```
# out: 1 x 1 x 1
nn.Flatten(),
nn.Sigmoid()
```

Note that we're using the Leaky ReLU activation for the discriminator.



Different from the regular ReLU function, Leaky ReLU allows the pass of a small gradient signal for negative values. As a result, it makes the gradients from the discriminator flows stronger into the generator. Instead of passing a gradient (slope) of 0 in the back-prop pass, it passes a small negative gradient. - [Source](#)

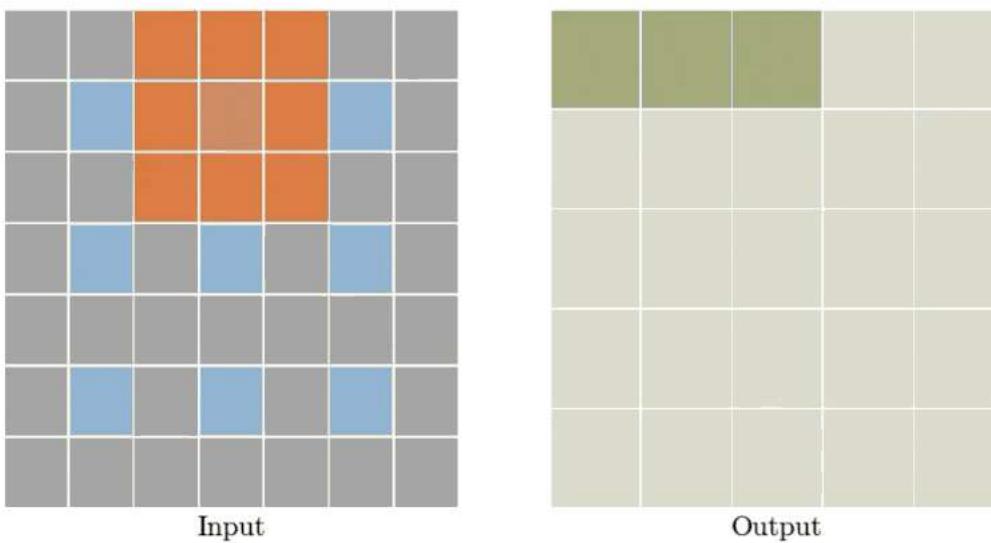
Just like any other binary classification model, the output of the discriminator is a single number between 0 and 1, which can be interpreted as the probability of the input image being real i.e. picked from the original dataset.

Let's move the discriminator model to the chosen device.

```
discriminator = to_device(discriminator, device)
```

▼ Generator Network

The input to the generator is typically a vector or a matrix of random numbers (referred to as a latent tensor) which is used as a seed for generating an image. The generator will convert a latent tensor of shape (128, 1, 1) into an image tensor of shape 3 x 28 x 28. To achieve this, we'll use the `ConvTranspose2d` layer from PyTorch, which performs a *transposed convolution* (also referred to as a *deconvolution*). [Learn more](#)



```
latent_size = 128
```

```
generator = nn.Sequential(
    # in: latent_size x 1 x 1
    nn.ConvTranspose2d(latent_size, 512, kernel_size=4, stride=1, padding=0, bias=False),
    nn.BatchNorm2d(512),
    nn.ReLU(True),
    # out: 512 x 4 x 4

    nn.ConvTranspose2d(512, 256, kernel_size=4, stride=2, padding=1, bias=False),
    nn.BatchNorm2d(256),
    nn.ReLU(True),
    # out: 256 x 8 x 8
```

```

nn.ConvTranspose2d(256, 128, kernel_size=4, stride=2, padding=1, bias=False),
nn.BatchNorm2d(128),
nn.ReLU(True),
# out: 128 x 16 x 16

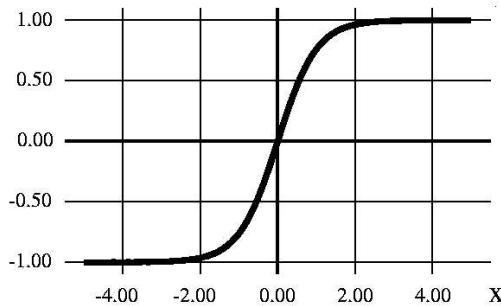
nn.ConvTranspose2d(128, 64, kernel_size=4, stride=2, padding=1, bias=False),
nn.BatchNorm2d(64),
nn.ReLU(True),
# out: 64 x 32 x 32

nn.ConvTranspose2d(64, 3, kernel_size=4, stride=2, padding=1, bias=False),
nn.Tanh()
# out: 3 x 64 x 64
)

```

We use the TanH activation function for the output layer of the generator.

hyperbolic tangent function
 $\tanh(x)$



"The ReLU activation (Nair & Hinton, 2010) is used in the generator with the exception of the output layer which uses the Tanh function. We observed that using a bounded activation allowed the model to learn more quickly to saturate and cover the color space of the training distribution. Within the discriminator we found the leaky rectified activation (Maas et al., 2013) (Xu et al., 2015) to work well, especially for higher resolution modeling." - [Source](#)

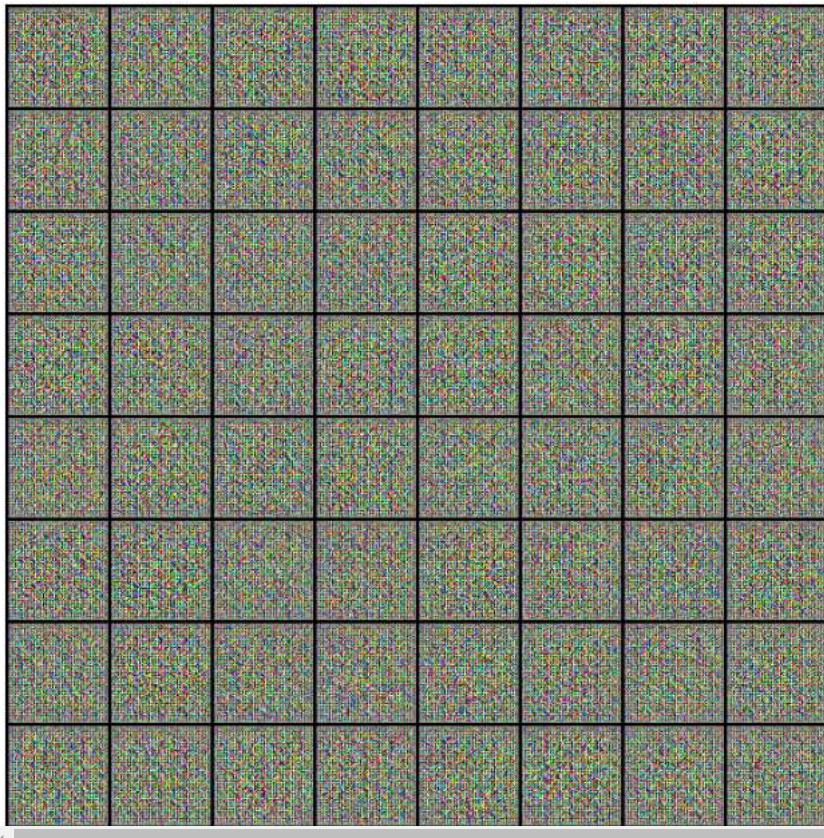
Note that since the outputs of the TanH activation lie in the range $[-1, 1]$, we have applied the similar transformation to the images in the training dataset. Let's generate some outputs using the generator and view them as images by transforming and denormalizing the output.

```

xb = torch.randn(batch_size, latent_size, 1, 1) # random latent tensors
fake_images = generator(xb)
print(fake_images.shape)
show_images(fake_images)

```

↳ `torch.Size([128, 3, 64, 64])`



```
generator = to_device(generator, device)
```

▼ Discriminator Training

Since the discriminator is a binary classification model, we can use the binary cross entropy loss function to quantify how well it is able to differentiate between real and generated images.

```
def train_discriminator(real_images, opt_d):
    # Clear discriminator gradients
    opt_d.zero_grad()

    # Pass real images through discriminator
    real_preds = discriminator(real_images)
    real_targets = torch.ones(real_images.size(0), 1, device=device)
    real_loss = F.binary_cross_entropy(real_preds, real_targets)
    real_score = torch.mean(real_preds).item()

    # Generate fake images
    latent = torch.randn(batch_size, latent_size, 1, 1, device=device)
    fake_images = generator(latent)

    # Pass fake images through discriminator
    fake_targets = torch.zeros(fake_images.size(0), 1, device=device)
    fake_preds = discriminator(fake_images)
    fake_loss = F.binary_cross_entropy(fake_preds, fake_targets)
    fake_score = torch.mean(fake_preds).item()

    # Update discriminator weights
    loss = real_loss + fake_loss
    loss.backward()
    opt_d.step()

    return loss.item(), real_score, fake_score
```

Here are the steps involved in training the discriminator.

- We expect the discriminator to output 1 if the image was picked from the real MNIST dataset, and 0 if it was generated using the generator network.
- We first pass a batch of real images, and compute the loss, setting the target labels to 1.
- Then we pass a batch of fake images (generated using the generator) pass them into the discriminator, and compute the loss, setting the target labels to 0.

- Finally we add the two losses and use the overall loss to perform gradient descent to adjust the weights of the discriminator.

It's important to note that we don't change the weights of the generator model while training the discriminator (`opt_d` only affects the `discriminator.parameters()`)

Generator Training

Since the outputs of the generator are images, it's not obvious how we can train the generator. This is where we employ a rather elegant trick, which is to use the discriminator as a part of the loss function. Here's how it works:

- We generate a batch of images using the generator, pass them into the discriminator.
- We calculate the loss by setting the target labels to 1 i.e. real. We do this because the generator's objective is to "fool" the discriminator.
- We use the loss to perform gradient descent i.e. change the weights of the generator, so it gets better at generating real-like images to "fool" the discriminator.

Here's what this looks like in code.

```
def train_generator(opt_g):
    # Clear generator gradients
    opt_g.zero_grad()

    # Generate fake images
    latent = torch.randn(batch_size, latent_size, 1, 1, device=device)
    fake_images = generator(latent)

    # Try to fool the discriminator
    preds = discriminator(fake_images)
    targets = torch.ones(batch_size, 1, device=device)
    loss = F.binary_cross_entropy(preds, targets)

    # Update generator weights
    loss.backward()
    opt_g.step()

    return loss.item()
```

Let's create a directory where we can save intermediate outputs from the generator to visually inspect the progress of the model. We'll also create a helper function to export the generated images.

```
from torchvision.utils import save_image

sample_dir = 'generated'
os.makedirs(sample_dir, exist_ok=True)

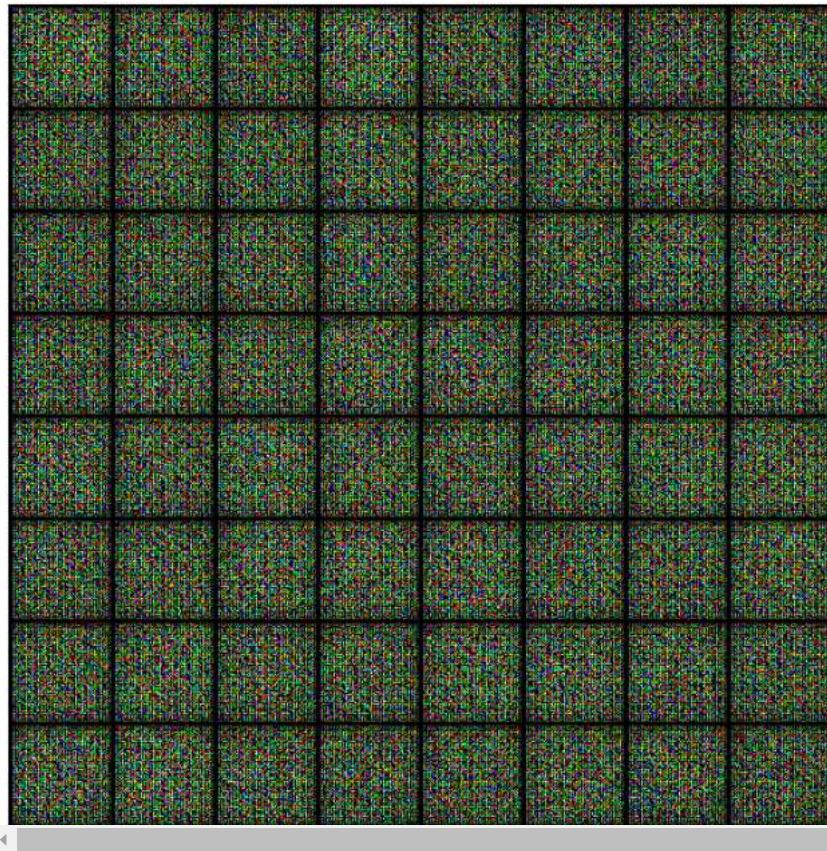
def save_samples(index, latent_tensors, show=False):
    fake_images = generator(latent_tensors)
    fake_fname = 'generated-images-{0:0=4d}.png'.format(index)
    save_image(denorm(fake_images), os.path.join(sample_dir, fake_fname), nrow=8)
    print('Saving', fake_fname)
    if show:
        fig, ax = plt.subplots(figsize=(8, 8))
        ax.set_xticks([]); ax.set_yticks([])
        ax.imshow(make_grid(fake_images.cpu().detach(), nrow=8).permute(1, 2, 0))
```

We'll use a fixed set of input vectors to the generator to see how the individual generated images evolve over time as we train the model. Let's save one set of images before we start training our model.

```
fixed_latent = torch.randn(64, latent_size, 1, 1, device=device)

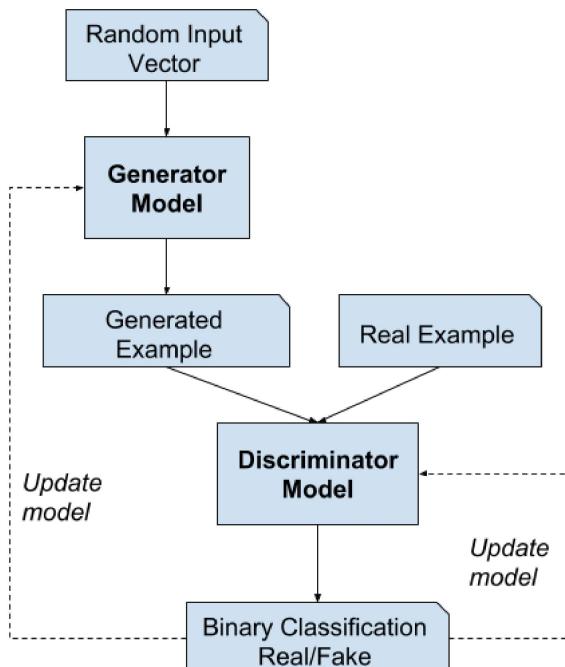
save_samples(0, fixed_latent)
```

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers)
Saving generated-images-0000.png



Full Training Loop

Let's define a `fit` function to train the discriminator and generator in tandem for each batch of training data. We'll use the Adam optimizer with some custom parameters (betas) that are known to work well for GANs. We will also save some sample generated images at regular intervals for inspection.



```
from tqdm.notebook import tqdm
import torch.nn.functional as F
```

```
def fit(epochs, lr, start_idx=1):
    torch.cuda.empty_cache()

    # Losses & scores
```

```
losses_g = []
losses_d = []
real_scores = []
fake_scores = []

# Create optimizers
opt_d = torch.optim.Adam(discriminator.parameters(), lr=lr, betas=(0.5, 0.999))
opt_g = torch.optim.Adam(generator.parameters(), lr=lr, betas=(0.5, 0.999))

for epoch in range(epochs):
    for real_images, _ in tdm(train_dl):
        # Train discriminator
        loss_d, real_score, fake_score = train_discriminator(real_images, opt_d)
        # Train generator
        loss_g = train_generator(opt_g)

    # Record losses & scores
    losses_g.append(loss_g)
    losses_d.append(loss_d)
    real_scores.append(real_score)
    fake_scores.append(fake_score)

    # Log losses & scores (last batch)
    print("Epoch [{}/{}], loss_g: {:.4f}, loss_d: {:.4f}, real_score: {:.4f}, fake_score: {:.4f}".format(
        epoch+1, epochs, loss_g, loss_d, real_score, fake_score))

    # Save generated images
    save_samples(epoch+start_idx, fixed_latent, show=False)

return losses_g, losses_d, real_scores, fake_scores
```

We are now ready to train the model. Try different learning rates to see if you can maintain the fine balance between the training the generator and the discriminator.

```
lr = 0.0002
epochs = 30

history = fit(epochs, lr)
```

```

[ 0% | 0/497 [00:00<?, ?it/s]
Epoch [1/30], loss_g: 3.9533, loss_d: 0.4527, real_score: 0.8454, fake_score: 0.2134
Saving generated-images-0001.png
 0% | 0/497 [00:00<?, ?it/s]
Epoch [2/30], loss_g: 4.6824, loss_d: 0.4092, real_score: 0.8614, fake_score: 0.2039
Saving generated-images-0002.png
 0% | 0/497 [00:00<?, ?it/s]
Epoch [3/30], loss_g: 4.4018, loss_d: 0.3086, real_score: 0.8073, fake_score: 0.0449
Saving generated-images-0003.png
 0% | 0/497 [00:00<?, ?it/s]
Epoch [4/30], loss_g: 4.8576, loss_d: 0.2428, real_score: 0.8629, fake_score: 0.0487
Saving generated-images-0004.png
 0% | 0/497 [00:00<?, ?it/s]
Epoch [5/30], loss_g: 5.7984, loss_d: 0.1232, real_score: 0.9530, fake_score: 0.0676
Saving generated-images-0005.png
 0% | 0/497 [00:00<?, ?it/s]
Epoch [6/30], loss_g: 5.9624, loss_d: 0.2209, real_score: 0.9092, fake_score: 0.0869
Saving generated-images-0006.png
 0% | 0/497 [00:00<?, ?it/s]
Epoch [7/30], loss_g: 6.3088, loss_d: 0.1087, real_score: 0.9665, fake_score: 0.0620
Saving generated-images-0007.png
 0% | 0/497 [00:00<?, ?it/s]
Epoch [8/30], loss_g: 5.4103, loss_d: 0.1025, real_score: 0.9335, fake_score: 0.0192
Saving generated-images-0008.png
 0% | 0/497 [00:00<?, ?it/s]
Epoch [9/30], loss_g: 6.6498, loss_d: 0.1475, real_score: 0.9307, fake_score: 0.0546
Saving generated-images-0009.png
 0% | 0/497 [00:00<?, ?it/s]
Epoch [10/30], loss_g: 9.3625, loss_d: 0.0281, real_score: 0.9813, fake_score: 0.0002
Saving generated-images-0010.png
 0% | 0/497 [00:00<?, ?it/s]
Epoch [11/30], loss_g: 5.6452, loss_d: 0.0974, real_score: 0.9523, fake_score: 0.0318
Saving generated-images-0011.png
 0% | 0/497 [00:00<?, ?it/s]
Epoch [12/30], loss_g: 8.4490, loss_d: 0.1164, real_score: 0.9536, fake_score: 0.0486
Saving generated-images-0012.png
 0% | 0/497 [00:00<?, ?it/s]
Epoch [13/30], loss_g: 14.9422, loss_d: 0.0767, real_score: 0.9934, fake_score: 0.0585
Saving generated-images-0013.png
 0% | 0/497 [00:00<?, ?it/s]
Epoch [14/30], loss_g: 7.4359, loss_d: 0.1557, real_score: 0.9663, fake_score: 0.0948
Saving generated-images-0014.png
 0% | 0/497 [00:00<?, ?it/s]
Epoch [15/30], loss_g: 7.2165, loss_d: 0.1051, real_score: 0.9638, fake_score: 0.0476
Saving generated-images-0015.png
 0% | 0/497 [00:00<?, ?it/s]
Epoch [16/30], loss_g: 7.8867, loss_d: 0.1831, real_score: 0.9138, fake_score: 0.0497
Saving generated-images-0016.png
 0% | 0/497 [00:00<?, ?it/s]
Epoch [17/30], loss_g: 18.5801, loss_d: 0.0334, real_score: 0.9924, fake_score: 0.0246
Saving generated-images-0017.png
 0% | 0/497 [00:00<?, ?it/s]
Epoch [18/30], loss_g: 9.1356, loss_d: 0.0382, real_score: 0.9994, fake_score: 0.0356
Saving generated-images-0018.png
 0% | 0/497 [00:00<?, ?it/s]
Epoch [19/30], loss_g: 6.4938, loss_d: 0.0663, real_score: 0.9996, fake_score: 0.0616
Saving generated-images-0019.png
```

```

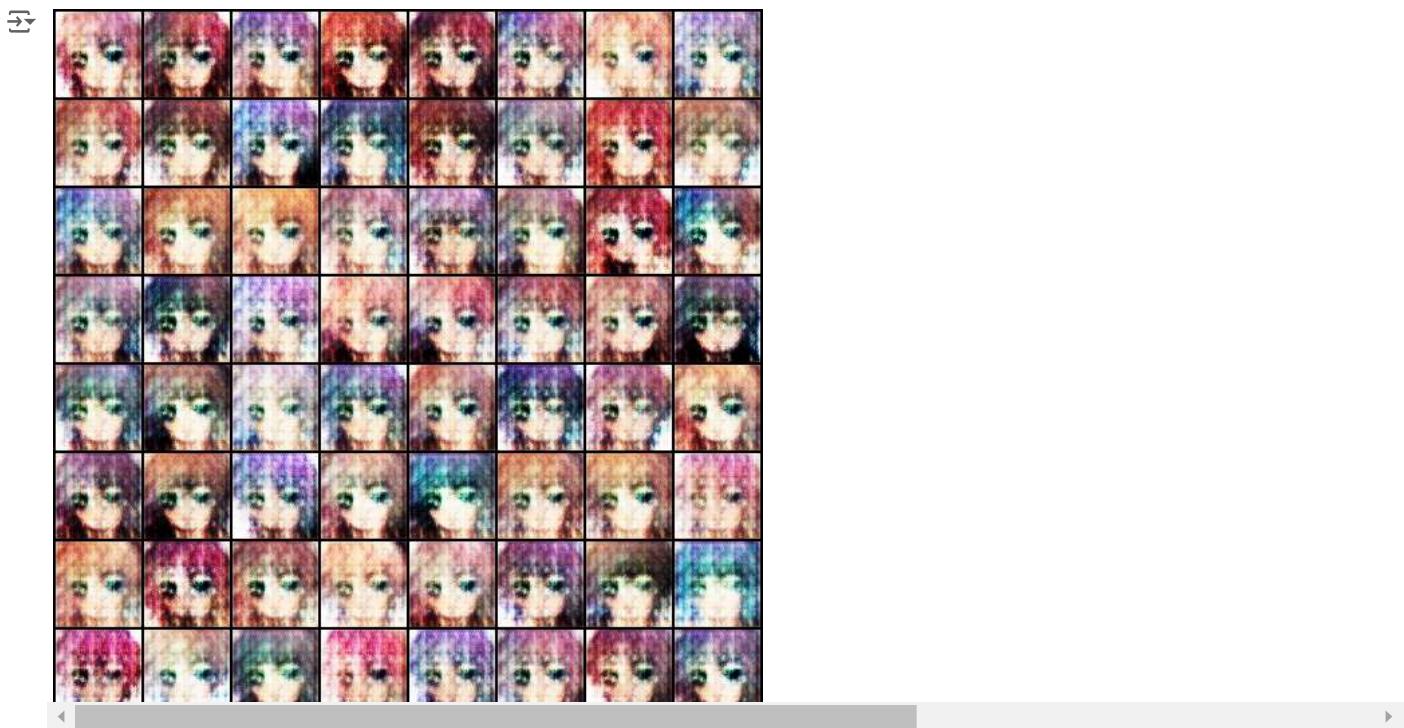
losses_g, losses_d, real_scores, fake_scores = history
Saving generated-images-0020.png
```

[link text](#) Now that we have trained the models, we can save checkpoints.

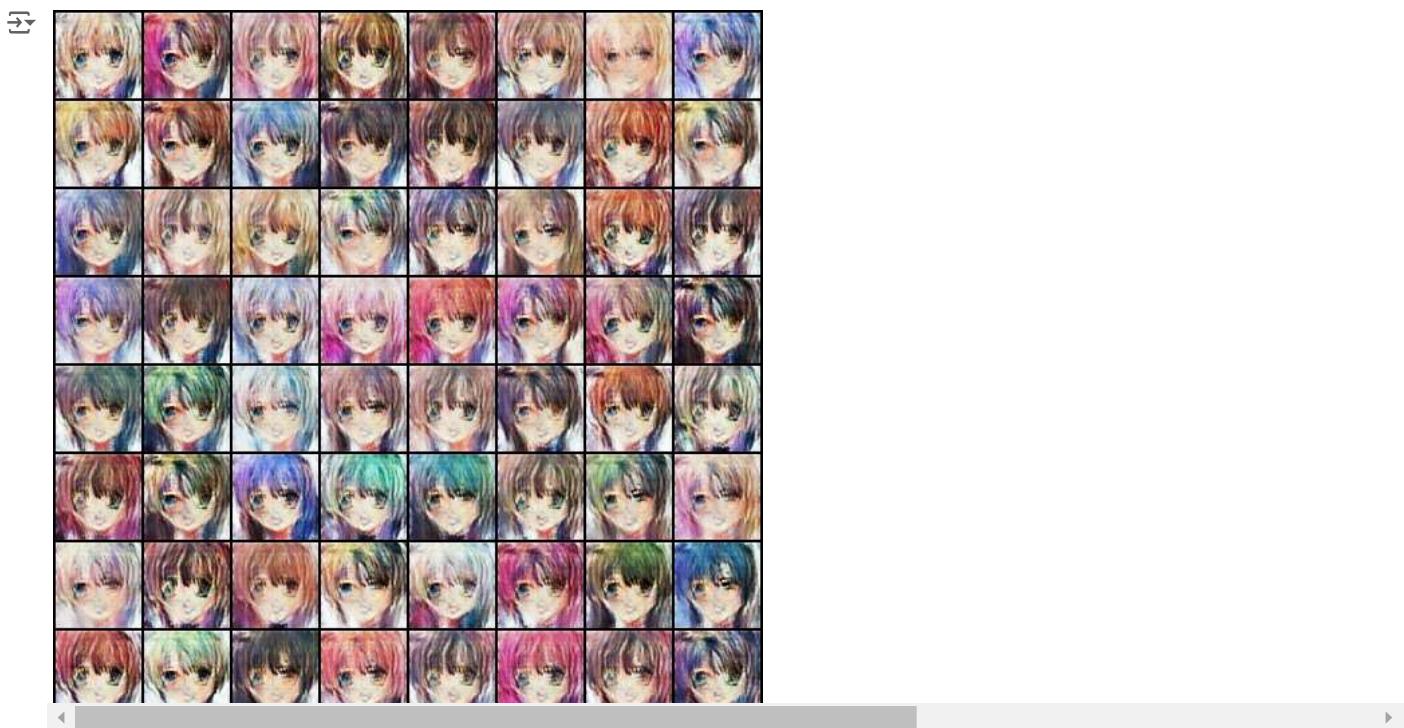
```
# Save the model checkpoints
torch.save(generator.state_dict(), 'G.pth')
torch.save(discriminator.state_dict(), 'D.pth')
```

Here's how the generated images look, after the 1st, 5th and 10th epochs of training.

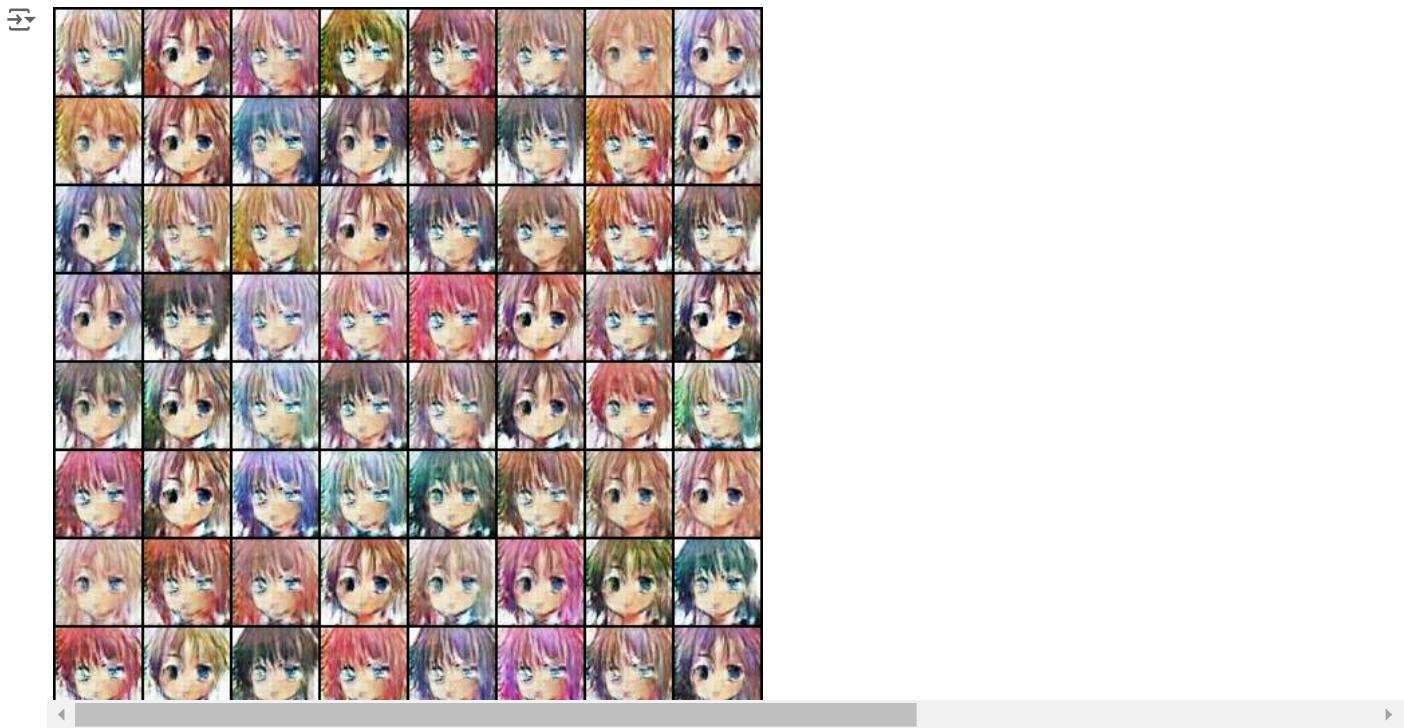
```
Saving generated-images-0024.png
from IPython.display import Image
Saving generated-images-0025.png
Image('./generated/generated-images-0001.png')
```



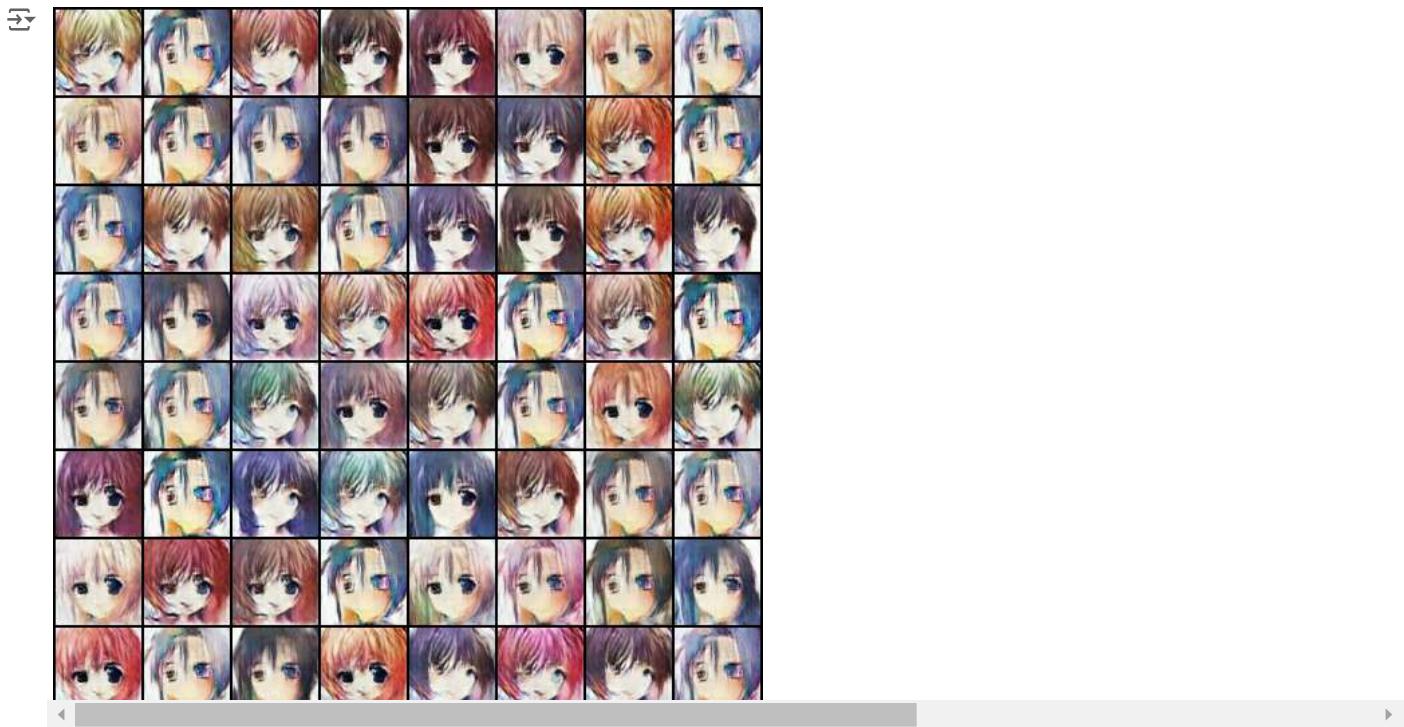
```
Image('./generated/generated-images-0005.png')
```



```
Image('./generated/generated-images-0010.png')
```



```
Image('./generated/generated-images-0020.png')
```



```
Image('./generated/generated-images-0025.png')
```

